

untangle::executor

Generated by Doxygen 1.18.0

1 executor	1
2 Namespace Index	7
2.1 Namespace List	7
3 Class Index	9
3.1 Class List	9
4 File Index	11
4.1 File List	11
5 Namespace Documentation	13
5.1 untangle Namespace Reference	13
5.1.1 Detailed Description	13
6 Class Documentation	15
6.1 untangle::executor< actionT > Class Template Reference	15
6.1.1 Detailed Description	15
6.1.2 Constructor & Destructor Documentation	16
6.1.2.1 executor()	16
6.1.2.2 ~executor()	16
6.1.3 Member Function Documentation	17
6.1.3.1 add_action()	17
6.1.3.2 add_task()	18
6.1.3.3 start()	19
6.1.3.4 stop()	19
6.1.4 Member Data Documentation	19
6.1.4.1 on_task_error	19
7 File Documentation	21
7.1 executor.hpp	21
Index	25

Chapter 1

executor

c++ thread pool executor based on async execution

The interface is one header that exposes a pool of workers, each one an **async** execution in continuous mode, fed by a shared queue of tasks.

This is the prototype, imported. It was written in `async/prototypes/` to answer whether execution is a usable building block for a pool — it is — and it is being hardened here. The behaviour below is what the tests state today, not a settled interface.

requirements

C++23, and the `async` submodule:

```
git submodule update --init --recursive
```

`async` carries its own actuator submodule, which is why the clone has to be recursive. Both are header-only; nothing is built or installed.

The headers report warnings with `std::println`, so they need a standard library that provides `<print>` — GCC 14 or newer.

example

```
#include <executor.hpp>
#include <functional>

int main() {
    untangle::executor<std::function<void(void)>> pool(4);
    pool.start();

    for (int i = 0; i < 100; ++i) {
        pool.add_action([i] { work(i); });
    }

    return 0; // ~executor finishes what was added, then stops the workers
}
```

the action type

The pool is a template on the callable it runs, the way an execution is — and on the same parameter name, because both kinds of queued work are that one callable. One pool runs one signature, and the caller names it:

```
untangle::executor<std::function<int(void)>> pool(4); // returns are run and dropped
```

The action type does not have to be a `std::function`. What an execution requires of it is a nested `result_type`, which it reads rather than deduces, so any callable declaring one will do — and a pool built on such a callable never touches `std::function::result_type`, which the standard removed in C++20 and both `libc++` and `libstdc++` still supply by grace rather than contract.

```
struct my_task {
    using result_type = void;
    void operator()() const;
};
```

```
untangle::executor<my_task> pool(4);
```

An action type may take arguments, and `add_action()` binds them the way its namesake `execution::add_action()` does — at the door, where the caller still holds them:

```
untangle::executor<std::function<int(int)>> pool(4);
pool.start();
pool.add_action([](int n) { return n * 2; }, 21);
```

What waits for a worker is one callable carrying its own arguments, so nothing about a task changes between the queue and the worker that runs it. The arguments are copied at `add_action()` and handed to the task as the pool's own lvalues when it runs: a task is called later, possibly much later, and what the caller passed may be gone by then. A task taking `int&` therefore mutates the pool's copy, which the caller never sees.

A task that cannot be called with the arguments given is refused by a `static_assert` in `add_action()`, rather than by an error from inside the binding.

tasks: work that reports back

Everything above queues work through `add_action()`, which is fire and forget: it runs, what it returns is dropped, and the caller hears nothing more about it. `add_task()` carries the callback it must notify:

```
untangle::executor<std::function<int(int)>> pool(4);
pool.start();

// the callback comes last, after the arguments the task is bound to
pool.add_task([](int n) { return n * 2; }, 21, [](int result) { /* result == 42 */ });
```

Work that returns void still reports that it finished, with a callback taking nothing — finished is the message and the result is optional:

```
untangle::executor<std::function<void(void)>> pool(4);
pool.add_task([] { /* ... */ }, [] { /* finished */ });
```

The callback is the last argument, and the signature cannot say so. A parameter pack cannot be followed by a deducible parameter, so it arrives inside the arguments and `untangle::bind_task()` splits it off. It must be callable with the task's result and return nothing — or callable with nothing at all, when the task returns void. A missing or unusable one is a compile error rather than silence.

Both doors share one queue, and the order out is the order in. An action posted before a task runs before it. The pool can promise that because it owns one container; `async::execution`, which the workers are, fires every action in a pass before any task, so its ordering across the two kinds differs. Only the pool's own order is a promise to a caller of the pool.

`pending()` counts both kinds, as it always counted work waiting for a worker.

A refused task does not notify. `add_task()` answers false under exactly the conditions `add_action()` does — before `start()`, after `stop()`, once the destructor has begun, or if the worker it was offered to has stopped — and a refused task is destroyed rather than queued. The answer is the whole report, so a caller waiting on the callback rather than on the answer would wait for ever. The same is true of work still queued when a stopped pool is destroyed: it never runs, so it never notifies, and the count on `stderr` is all there is.

Finished does not mean failed. A task that throws is not notified: there is no result to hand over, and for a void task no completion to report either. What it threw goes to `on_task_error`, as a failing action's does. And because a callback runs inside the same try as the task that owns it, `on_task_error` can fire for a task that in fact succeeded — the work was done and only the telling failed.

A callback runs on a worker's thread, inside that worker's drain, and several may be running at once on different workers. It is the same rule `on_task_error` carries, for the same reason: nothing may escape it.

how work is placed

N executions run in continuous mode behind a shared `std::deque` of tasks.

Nothing overtakes. A task goes straight to a free worker only when the shared queue is empty; otherwise it joins the back of the queue. Finding a free worker does not let a task jump a queue that already has work in it.

Workers pull rather than being pushed to. When one reports its queue finished — `on_finished`, raised by execution the moment its list empties — it takes the front of the shared queue. That is what spreads the work, with no scheduling logic of its own: 400 tasks over 4 workers come out 100/100/100/100.

A worker is asked, not tracked. `execution::is_busy()` is what `add_action()` reads to find a free worker and what the destructor reads to decide the pool is idle. A tally kept here would be the pool's belief about its workers; this is the workers' own answer, and it cannot drift.

running, stopping, restarting

A pool is built stopped. The constructor makes the workers and wires them up; `start()` runs them. Until it is called the pool refuses work, because a task taken then would sit unrun and the destructor would wait for it. It is also the moment to assign `on_task_error`, before anything can throw.

`stop()` stops the workers, and stops the pool accepting. What a worker is already running it finishes; what is still queued stays queued. `stop()` does not wait for the workers to leave their threads — only the destructor does that, and that wait is what makes destroying the pool safe.

`start()` after `stop()` restarts it, workers and all, and the queue it was holding drains as they pick it up again. Either call is harmless twice over.

lifetime and failure

The destructor finishes what was added. It stops accepting, waits for the queue to empty and for every worker to go idle, and only then stops them. A task added before the pool goes out of scope has run by the time it does.

A pool destroyed while stopped is the exception: its queue can never empty, because only a running worker takes from it, so the destructor does not wait for it. Whatever is left is reported on stderr and dropped.

Both waits are unbounded, and both say so while they last. A task that never returns would otherwise leave the destructor silent for ever, so it reports instead — the first wait naming who is still inside a task, the second naming whose thread has not left:

```
executor: still waiting to finish after 1s - 1 queued, 2 in a task: pool_worker_0, pool_worker_1
executor: still waiting to finish after 3s - 0 queued, 1 in a task: pool_worker_1
executor: still waiting to stop after 1s - 1 not left its thread: pool_worker_1
```

The interval starts at a second and doubles to a ceiling of thirty, so a stuck teardown says something almost at once without a long one becoming a flood. The two waits are kept apart because they answer different questions: a worker can be idle and still be in its thread.

Bounding either wait was never available. Returning early while a worker is still in a task is a use-after-free, and a bounded wait only moves the hang into `~execution()`, which spins on its own running state with no timeout. Workers are detached and `async::execution` has no cancellation, so a pool cannot outlive a task it cannot interrupt. What can be fixed is the silence.

`add_action()` says whether the task was taken. It returns false before `start()`, after `stop()`, and once the destructor has begun. It returns false again if the worker it was offered to has stopped. A refused task is destroyed rather than queued, so a false answer means it will not run.

A task that is still running is refused too, and that is chosen rather than incidental. A pool shutting down is waiting for exactly that task, so it is tempting to let the work it spawns through — but then a task that re-adds itself could keep the shutdown from ever ending. One rule instead: a pool that is not running takes no work, whoever is asking. `add_action()` answers false inside the task, so it can tell.

A task that throws does not take the worker with it, and what it threw is not lost. The worker catches it and carries on with the next task; `on_task_error` is where the throw goes:

```
untangle::executor<std::function<void(void)>> pool(4);

pool.on_task_error = [](std::exception_ptr thrown) {
    try {
        std::rethrow_exception(thrown);
    } catch (const std::exception& e) {
        log(e.what());
    }
};
```

It carries a `std::exception_ptr` because a task may throw something that is not a `std::exception`, and that is the case most worth hearing about. Assign it before the first `add_action()` — a worker reads it — and expect it on a worker's thread, with more than one worker possibly inside it at once. Leave it unset and the pool warns on stderr instead, naming the worker.

That name is not unique between pools. Workers are named by index — `pool_worker_0`, `pool_worker_1` — so two pools in one process each have a `pool_worker_0`, and every warning that names one is ambiguous: the failure reports above, and the destructor's waits alike. With a single pool the name says which worker; with two it does not. Assigning `on_task_error` sidesteps it for failures, since the handler is the pool's own and the warning is replaced.

A return value is still not collected: a continuous worker keeps none. Work with something to report carries its own channel, which is what `add_task()` is for — see tasks: work that reports back above.

building the tests

```
cmake -S test -B test/build
cmake --build test/build
ctest --test-dir test/build
```

test/CMakeLists.txt fetches googletest at configure time, so the first configure needs network access.

sanitizers

EXECUTOR_SANITIZE builds the tests under a sanitizer. It is off by default, because a sanitized build is several times slower and ThreadSanitizer does not ship for every toolchain.

```
cmake -S test -B test/build-asan -DEXECUTOR_SANITIZE=address
cmake --build test/build-asan
ctest --test-dir test/build-asan
```

Accepted values are address, thread, undefined, or empty. Anything else is refused at configure time rather than passed through to the compiler.

Use a separate build directory per sanitizer. address and thread instrument the same accesses in incompatible ways and cannot be combined, so one build directory is one sanitizer.

the reference

tools/make_doc.sh builds doc/refman.pdf from the header and this README. It needs doxygen, tectonic and python3, and it fails on any doxygen warning that is not an obsolete-tag notice, so an undocumented member or a broken \ref does not reach the reference.

Chapter 2

Namespace Index

2.1 Namespace List

Here is a list of all documented namespaces with brief descriptions:

untangle	A thread pool executor built out of <code>untangle::async::execution</code> objects	13
--------------------------	---	--------------------

Chapter 3

Class Index

3.1 Class List

Here are the classes, structs, unions and interfaces with brief descriptions:

[untangle::executor< actionT >](#)
Runs actions and tasks on a fixed pool of untangle::async::execution workers [15](#)

Chapter 4

File Index

4.1 File List

Here is a list of all documented files with brief descriptions:

executor.hpp	21
--	----

Chapter 5

Namespace Documentation

5.1 untangle Namespace Reference

A thread pool executor built out of `untangle::async::execution` objects.

Classes

- class `executor`
Runs actions and tasks on a fixed pool of `untangle::async::execution` workers.

5.1.1 Detailed Description

A thread pool executor built out of `untangle::async::execution` objects.

N workers run in continuous mode behind a shared queue. A task goes straight to a free worker only when the queue is empty; otherwise it joins the back, and a worker takes the front as it frees up.

Chapter 6

Class Documentation

6.1 `untangle::executor< actionT >` Class Template Reference

Runs actions and tasks on a fixed pool of `untangle::async::execution` workers.

```
#include <executor.hpp>
```

Public Member Functions

- `executor` (`std::size_t worker_count`)
Builds `worker_count` workers. They do not run until `start()`.
- `~executor` ()
Finishes what has been added, then stops every worker.
- `template<typename... Args>`
`bool add_action` (`actionT task`, `Args &&... args`)
Queues a task bound to `args`, or hands it to a worker that has nothing to do.
- `template<typename... Args>`
`bool add_task` (`actionT task`, `Args &&... args`)
Queues work that must report, or hands it to a worker that has nothing to do.
- `std::size_t pending` () `const`
How many tasks are waiting for a worker.
- `void start` ()
Starts every worker, and starts the pool accepting work.
- `void stop` ()
Stops every worker. What they are already running, they finish.

Public Attributes

- `std::function< void(std::exception_ptr)>` `on_task_error`
Called with whatever the work threw. Assigned by the caller.

6.1.1 Detailed Description

```
template<typename actionT>  
class untangle::executor< actionT >
```

Runs actions and tasks on a fixed pool of `untangle::async::execution` workers.

Template Parameters

actionT	The callable the pool runs. One pool runs one signature - its arguments as much as its result.
---------	--

Remarks

****Two kinds of work may be queued, and both are an actionT.**** [add_action\(\)](#) is fire and forget: the work runs, what it returns is dropped, and nothing reports that it finished. [add_task\(\)](#) carries the callback it must notify, and reports what the work produced - or, for a callable returning void, merely that it finished. The names and the split are `async::execution`'s, which this is built out of.

6.1.2 Constructor & Destructor Documentation

6.1.2.1 executor()

```
template<typename actionT>
untangle::executor< actionT >::executor (
    std::size_t worker_count) [inline], [explicit]
```

Builds `worker_count` workers. They do not run until [start\(\)](#).

Parameters

<code>worker_count</code>	- How many workers to run. Must not be zero.
---------------------------	--

Exceptions

<code>std::invalid_argument</code>	- <code>worker_count</code> is zero. Such a pool would take work it could never run.
------------------------------------	--

6.1.2.2 ~executor()

```
template<typename actionT>
untangle::executor< actionT >::~~executor () [inline]
```

Finishes what has been added, then stops every worker.

Both waits are unbounded, and both report on `stderr` while they last, naming the workers they are waiting on.

Attention

A task still running here is refused if it adds more work, even though this is waiting for that very task. It is deliberate: a shutdown that took new work could be kept from ever ending by a task that re-adds itself. [add_action\(\)](#) and [add_task\(\)](#) answer false, so the task can see it.

Queued work dropped here never runs, so a task dropped here never notifies. The count on `stderr` is the whole report, and a caller waiting on a callback rather than on the answer would wait for ever. Destroying a pool with work still queued is how that happens; the destructor says how much was lost but not which.

6.1.3 Member Function Documentation

6.1.3.1 add_action()

```
template<typename actionT>
template<typename... Args>
bool untangle::executor< actionT >::add_action (
    actionT task,
    Args &&... args) [inline]
```

Queues a task bound to args, or hands it to a worker that has nothing to do.

Remarks

Named for `async::execution::add_action()`, which it forwards to and which behaves the same way: the work runs, what it returns is dropped, and nothing reports that it finished. The answer below is the last the caller hears.

The arguments are bound here, while the caller still holds them, and what waits for a worker is one callable carrying its own. That is the shape `async::execution::add_action()` offers at its door, and it is bound here for the same reason it is bound there: a queue holds callables, and a call site cannot be queued.

Remarks

The first free worker takes it, counting from the start, so a pool that is not busy keeps giving work to the same one. That is the intent: the worker that just ran a task is the warm one, and a free worker is free whichever it is.

Template Parameters

Args	- The argument types to bind to task.
------	---------------------------------------

Parameters

task	- The task to run.
args	- What to bind to it. Copied here and handed to the task as the pool's own lvalues when it runs, so a task taking a reference is given that copy rather than the caller's object, which may be gone by then.

Returns

true - the task was accepted. false - it was refused, by a pool that has not been started, has been stopped, or is being destroyed, or by a worker that has stopped. Either way it will not run. A task calling this from a worker thread is answered like any other caller.

6.1.3.2 add_task()

```
template<typename actionT>
template<typename... Args>
bool untangle::executor< actionT >::add_task (
    actionT task,
    Args &&... args) [inline]
```

Queues work that must report, or hands it to a worker that has nothing to do.

The task runs exactly as one given to [add_action\(\)](#) does, and then the callback it was built with is invoked with what it returned - on the worker's thread, inside the worker's drain.

Attention

The last argument is the callback, and this signature cannot say so: a parameter pack cannot be followed by a deducible parameter, so it arrives inside args and `untangle::bind_task()` splits it off. It must be callable with the task's result and return nothing - or callable with nothing at all, when the task returns void, because finished is the message and the result is optional. A missing or unusable one is a `static_assert` rather than silence.

Finished does not mean failed. A task that throws is not notified - there is no result to hand over - and what it threw goes to [on_task_error](#), as a failing action's does. A caller waiting on the callback alone would wait for ever.

A refused task does not notify either. The answer below is the whole report.

Remarks

It shares one queue with [add_action\(\)](#), and the order is the order they arrived: an action posted before a task runs before it.

Template Parameters

Args	- The argument types to bind, of which the last is the callback.
------	--

Parameters

task	- The task to run.
args	- What to bind to it, followed by the callback to notify. Copied here, as add_action() copies them, and for the same reason.

Returns

true - the task was accepted, and the callback will be notified when it has run.

false - refused, exactly as [add_action\(\)](#) is refused, and it will neither run nor notify.

6.1.3.3 start()

```
template<typename actionT>
void untangle::executor< actionT >::start () [inline]
```

Starts every worker, and starts the pool accepting work.

A pool refuses work until this is called, and refuses it again after [stop\(\)](#) until this is called a second time - a stopped pool is restarted by it, workers and all. Calling it twice over is harmless.

6.1.3.4 stop()

```
template<typename actionT>
void untangle::executor< actionT >::stop () [inline]
```

Stops every worker. What they are already running, they finish.

Queued tasks stay unrun and a task added afterwards is refused. That holds whoever is asking, a task already running on a worker included - a pool that is not running takes no work. [start\(\)](#) gives the workers a new working life. Calling it twice is harmless.

Attention

It does not wait for the workers to leave their threads. Only the destructor does that, and that wait is what makes destroying the pool safe.

6.1.4 Member Data Documentation

6.1.4.1 on_task_error

```
template<typename actionT>
std::function<void(std::exception_ptr)> untangle::executor< actionT >::on_task_error
```

Called with whatever the work threw. Assigned by the caller.

Nothing else reports failed work: [add_action\(\)](#) and [add_task\(\)](#) have both answered long before it runs. Left unset, the pool warns on stderr instead.

Remarks

It receives a `std::exception_ptr`, because the work may throw what is not a `std::exception`. Rethrow it to read it.

A task's failure arrives here as well, and so does its callback's. A task that throws is not notified - there is no result to hand over - so this is the only report it makes. And because a callback runs inside the same try as the task that owns it, this can fire for a task that in fact succeeded: the work was done and only the telling failed.

Attention

It runs on a worker's thread, and several workers may be in it at once. Assign it before the first [add_action\(\)](#) or [add_task\(\)](#), and let nothing escape it.

A task's callback runs on a worker's thread too, under the same rule and for the same reason - several tasks may be notifying at once, on different workers, and nothing may escape a callback either.

The documentation for this class was generated from the following file:

- `executor.hpp`

Chapter 7

File Documentation

7.1 executor.hpp

```
00001 // Copyright (c) 2026 Nicolae Popescu. MIT License.
00002
00010 #pragma once
00011
00012 #include <algorithm>
00013 #include <async.hpp>
00014 #include <atomic>
00015 #include <chrono>
00016 #include <condition_variable>
00017 #include <cstdint>
00018 #include <deque>
00019 #include <exception>
00020 #include <functional>
00021 #include <memory>
00022 #include <mutex>
00023 #include <print>
00024 #include <stdexcept>
00025 #include <string>
00026 #include <thread>
00027 #include <tuple>
00028 #include <type_traits>
00029 #include <utility>
00030 #include <vector>
00031
00032 namespace untangle {
00033
00046 template <typename actionT>
00047 class executor {
00048 public:
00057 explicit executor(std::size_t worker_count) {
00058     if (worker_count == 0) {
00059         throw std::invalid_argument("executor: worker_count must not be zero");
00060     }
00061
00062     workers_.reserve(worker_count);
00063
00064     for (std::size_t index = 0; index < worker_count; ++index) {
00065         auto next = worker_execution::create_instance("pool_worker_" + std::to_string(index));
00066
00067         // Tells the pool the moment this worker has room for more.
00068         next->on_finished = [this, index] { take_next_task(index); };
00069
00070         // Carries what a task throws to the caller. Assigned here because a worker thread reads it.
00071         next->on_error = [this, name = next->name](std::exception_ptr thrown) {
00072             report_task_error(name, thrown);
00073         };
00074
00075         // Adds the worker to the poll the destructor waits on.
00076         poll_.add(*next);
00077
00078         workers_.push_back(std::move(next));
00079     }
00080 }
00081
00098 ~executor() {
00099     // Whether the pool was still running when it was destroyed. A stopped one cannot empty its
00100     // queue - only a worker takes from it, through on_finished, and a stopped worker raises no more
```

```

00101 // - so its queue is not waited for.
00102 const bool was_running = running_.exchange(false);
00103
00104 {
00105     std::unique_lock<std::mutex> lock(mutex_);
00106
00107     // Waits for the queue to empty and every worker to go idle.
00108     auto interval = std::chrono::milliseconds(report_first_ms);
00109     auto waited = std::chrono::milliseconds(0);
00110
00111     while (!finished_cv_.wait_for(lock, interval, [this, was_running] {
00112         return nothing_running() && (pending_.empty() || !was_running);
00113     })) {
00114         waited += interval;
00115
00116         // Under the lock, so the counts agree with the predicate that just failed.
00117         std::println(
00118             stderr, "executor: still waiting to finish after {}s - {} queued, {} in a task{}",
00119             waited.count() / 1000, pending_.size(), count_workers(&worker_execution::is_busy),
00120             name_workers(&worker_execution::is_busy));
00121
00122         interval = std::min(interval * 2, std::chrono::milliseconds(report_max_ms));
00123     }
00124
00125     if (!pending_.empty()) {
00126         std::println(stderr,
00127             "executor: {} queued task(s) dropped, the pool was stopped before they ran",
00128             pending_.size());
00129     }
00130 }
00131
00132 stop();
00133
00134 // Waits for every worker to leave its thread, which is what makes destroying them safe. A
00135 // different question from the wait above: a worker can be idle and still be in its thread.
00136 auto interval = std::chrono::milliseconds(report_first_ms);
00137 auto waited = std::chrono::milliseconds(0);
00138
00139 while (poll_.is_running()) { // time of check
00140     const auto tick = std::chrono::milliseconds(poll_interval_ms);
00141     std::this_thread::sleep_for(tick);
00142     waited += tick;
00143
00144     if (waited >= interval) {
00145         std::println(stderr, "executor: still waiting to stop after {}s - {} not left its thread{}",
00146             waited.count() / 1000, count_workers(&worker_execution::is_running),
00147             name_workers(&worker_execution::is_running));
00148
00149         interval = std::min(interval * 2, std::chrono::milliseconds(report_max_ms));
00150     }
00151 }
00152
00153 // time of use: no worker thread is left to reach workers_.
00154 workers_.clear();
00155 }
00156
00157 template <typename... Args>
00158 bool add_action(actionT task, Args&&... args) {
00159     // The call the pool will make, which is not quite the one written at the call site: what
00160     // reaches the task is the copies bound below, so that is the call that has to compile.
00161     static_assert(std::is_invocable_v<actionT, std::decay_t<Args>&&...>,
00162         "executor: the task cannot be called with the arguments given to add_action()");
00163
00164     if (!running_) {
00165         return false;
00166     }
00167
00168     // Bound before the lock: it copies the caller's arguments, and nothing here needs the queue.
00169     return queue(bind_action(std::move(task), std::forward<Args>(args)...));
00170 }
00171
00172 template <typename... Args>
00173 bool add_task(actionT task, Args&&... args) {
00174     if (!running_) {
00175         return false;
00176     }
00177
00178     // untangle::bind_task() takes the callback off the end of the pack and checks it; what comes
00179     // back is one callable that runs the work and then notifies, which is what the queue holds.
00180     return queue(untangle::bind_task(std::move(task), std::forward<Args>(args)...));
00181 }
00182
00183 std::size_t pending() const {
00184     std::lock_guard<std::mutex> lock(mutex_);
00185     return pending_.size();
00186 }
00187
00188
00189
00190
00191
00192
00193
00194
00195
00196
00197
00198
00199
00200
00201
00202
00203
00204
00205
00206
00207
00208
00209
00210
00211
00212
00213
00214
00215
00216
00217
00218
00219
00220
00221
00222
00223
00224
00225
00226
00227
00228
00229
00230
00231
00232
00233
00234
00235
00236
00237
00238
00239
00240
00241
00242
00243
00244
00245
00246
00247
00248

```

```

00256 void start() {
00257     // The workers first, so anything accepted after this is handed to one that is running.
00258     for (auto& one : workers_) {
00259         one->start();
00260     }
00261
00262     running_ = true;
00263 }
00264
00275 void stop() {
00276     // First, so that nothing is accepted for workers that are about to stop. A task already on a
00277     // worker still runs; one arriving from here is refused.
00278     running_ = false;
00279
00280     // No lock: a worker waking from stop() runs on_finished, which takes mutex_.
00281     for (auto& one : workers_) {
00282         one->stop();
00283     }
00284 }
00285
00307 std::function<void(std::exception_ptr)> on_task_error;
00308
00309 private:
00316 void report_task_error(const std::string& worker, std::exception_ptr thrown) {
00317     if (on_task_error) {
00318         on_task_error(thrown);
00319         return;
00320     }
00321
00322     try {
00323         std::rethrow_exception(thrown);
00324     } catch (const std::exception& e) {
00325         std::println(stderr, "warning: executor task on '{}' threw: {}", worker, e.what());
00326     } catch (...) {
00327         std::println(stderr, "warning: executor task on '{}' threw an unknown type", worker);
00328     }
00329 }
00330
00341 template <typename... Args>
00342 static untangle::task_t bind_action(actionT task, Args&&... args) {
00343     if constexpr (sizeof...(Args) == 0) {
00344         return untangle::task_t(std::move(task));
00345     } else {
00346         // Captured by value and called as lvalues, which is all a call deferred to a worker can
00347         // promise: whatever was passed to add_action() may be gone by the time this runs.
00348         return untangle::task_t(
00349             [task = std::move(task), ... args = std::forward<Args>(args)]() mutable { task(args...); });
00350     }
00351 }
00352
00353 using worker_execution = untangle::async::execution<untangle::task_t>;
00354
00363 bool queue(untangle::task_t call) {
00364     std::lock_guard<std::mutex> lock(mutex_);
00365
00366     // The queue comes first: an entry does not overtake one already waiting in it.
00367     if (pending_.empty()) {
00368         // Always from the start, so a free pool reuses its warmest worker rather than waking a cold
00369         // one. Under load the workers are busy and the queue below is what spreads the work.
00370         for (std::size_t index = 0; index < workers_.size(); ++index) {
00371             if (!workers_[index]->is_busy()) {
00372                 // A refused entry is already destroyed, and stop() stops every worker together, so there
00373                 // is nothing to queue and no other worker to try.
00374                 return give_to_worker(index, std::move(call));
00375             }
00376         }
00377     }
00378
00379     pending_.push_back(std::move(call));
00380     return true;
00381 }
00382
00394 [[nodiscard]] bool give_to_worker(std::size_t index, untangle::task_t task) {
00395     // The door takes the worker's action_mutex under mutex_. The two never cycle, because a worker
00396     // calls on_finished with its action_mutex released.
00397     return workers_[index]->add_action(std::move(task));
00398 }
00399
00403 void take_next_task(std::size_t index) {
00404     std::lock_guard<std::mutex> lock(mutex_);
00405
00406     if (!pending_.empty()) {
00407         untangle::task_t next = std::move(pending_.front());
00408         pending_.pop_front();
00409
00410         if (!give_to_worker(index, std::move(next))) {
00411             // The task is already destroyed, so it is reported rather than put back.

```

```

00412         std::println(stderr, "executor: {} refused a queued task, which is lost",
00413                         workers_[index]->name);
00414     }
00415     return;
00416 }
00417
00418 if (nothing_running()) {
00419     finished_cv_.notify_all();
00420 }
00421 }
00422 }
00423
00424 bool nothing_running() const {
00425     for (const auto& one : workers_) {
00426         if (one->is_busy()) {
00427             return false;
00428         }
00429     }
00430     return true;
00431 }
00432
00433 std::size_t count_workers(bool (worker_execution::*state)() const) const {
00434     std::size_t count = 0;
00435
00436     for (const auto& one : workers_) {
00437         if ((one.get()->*state)()) {
00438             ++count;
00439         }
00440     }
00441     return count;
00442 }
00443
00444 std::string name_workers(bool (worker_execution::*state)() const) const {
00445     std::string names;
00446
00447     for (const auto& one : workers_) {
00448         if ((one.get()->*state)()) {
00449             names += names.empty() ? " : " : ", ";
00450             names += one->name;
00451         }
00452     }
00453     return names;
00454 }
00455
00456 static constexpr int poll_interval_ms = 50;
00457
00458 static constexpr int report_first_ms = 1000;
00459 static constexpr int report_max_ms = 30000;
00460
00461 mutable std::mutex mutex_;
00462 std::condition_variable finished_cv_;
00463
00464 async::execution_poll poll_;
00465
00466 std::deque<untangle::task_t> pending_;
00467 std::vector<std::shared_ptr<worker_execution>> workers_;
00468 std::atomic_bool running_ = {false};
00469 };
00470
00471 // namespace untangle

```

Index

- ~executor
 - untangle::executor< actionT >, [16](#)
- add_action
 - untangle::executor< actionT >, [17](#)
- add_task
 - untangle::executor< actionT >, [17](#)
- executor, [1](#)
 - untangle::executor< actionT >, [16](#)
- on_task_error
 - untangle::executor< actionT >, [19](#)
- start
 - untangle::executor< actionT >, [18](#)
- stop
 - untangle::executor< actionT >, [19](#)
- untangle, [13](#)
- untangle::executor< actionT >, [15](#)
 - ~executor, [16](#)
 - add_action, [17](#)
 - add_task, [17](#)
 - executor, [16](#)
 - on_task_error, [19](#)
 - start, [18](#)
 - stop, [19](#)